

# Data 412 HW 4

Ashley Totten

February 20, 2025

## Star Wars

```
library(dplyr)
```

Attaching package: 'dplyr'

The following objects are masked from 'package:stats':

filter, lag

The following objects are masked from 'package:base':

intersect, setdiff, setequal, union

```
data(starwars)
```

1. The only difference in the output is that the group\_by function put the droid species first in the tibble whereas the .by function put the human species first.

```
starwars %>%
  filter(homeworld == "Tatooine") %>%
  group_by(species) %>%
  summarize(height_mean = mean(height), height_sd = sd(height), mass_mean = mean(mass), mass_sd = sd(mass))
```

```
# A tibble: 2 x 6
  species height_mean height_sd mass_mean mass_sd   num
  <chr>        <dbl>     <dbl>     <dbl>    <dbl> <int>
1 Droid         132      49.5     53.5    30.4     2
2 Human        179.     12.8      NA       NA      8
```

```
starwars %>%
  filter(homeworld == "Tatooine") %>%
  summarize(.by = species, height_mean = mean(height), height_sd = sd(height), mass_mean = me
```

```
# A tibble: 2 x 6
  species height_mean height_sd mass_mean mass_sd   num
  <chr>        <dbl>     <dbl>     <dbl>    <dbl> <int>
1 Human        179.     12.8      NA       NA      8
2 Droid         132      49.5     53.5    30.4     2
```

2. The tallest droid is C-3PO. The tallest human is Darth Vader.

```
starwars %>%
  filter(homeworld == "Tatooine") %>%
  select(name, species, height) %>%
  arrange(species, desc(height))
```

```
# A tibble: 10 x 3
  name            species height
  <chr>           <chr>    <int>
1 C-3PO          Droid     167
2 R5-D4          Droid     97
3 Darth Vader    Human    202
4 Anakin Skywalker Human   188
5 Biggs Darklighter Human   183
6 Cliegg Lars   Human    183
7 Owen Lars     Human    178
8 Luke Skywalker Human   172
9 Beru Whitesun Lars Human   165
10 Shmi Skywalker Human   163
```

3. The smallest droid is R5-D4, he appeared in a New Hope. The smallest human is Beru Whitesun Lars who appeared in A New Hope, Attack of the Clones, and Revenge of the Sith.

```
starwars %>%
  filter(homeworld == "Tatooine") %>%
  select(name, species, mass, films) %>%
  arrange(species, mass)
```

```
# A tibble: 10 x 4
  name          species   mass films
  <chr>        <chr>     <dbl> <list>
1 R5-D4         Droid      32 <chr [1]>
2 C-3PO          Droid      75 <chr [6]>
3 Beru Whitesun Lars Human    75 <chr [3]>
4 Luke Skywalker Human     77 <chr [5]>
5 Biggs Darklighter Human    84 <chr [1]>
6 Anakin Skywalker Human    84 <chr [3]>
7 Owen Lars      Human   120 <chr [3]>
8 Darth Vader    Human   136 <chr [4]>
9 Shmi Skywalker Human    NA <chr [2]>
10 Cliegg Lars     Human   NA <chr [1]>
```

## NYC Flights

```
library(nycflights13)
data(flights)
```

1. There are 10200 flights that had an arrival delay of 2 hours or more. There are 139504 flights that were operated by United, American, or Delta.

```
flights %>%
  filter(arr_delay >= 120) -> delay_2
head(delay_2, 10)
```

```
# A tibble: 10 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>           <int>     <dbl>    <int>           <int>
1 2013     1     1      811            630       101      1047           830
2 2013     1     1      848            1835      853      1001          1950
3 2013     1     1      957            733      144      1056           853
4 2013     1     1     1114            900      134      1447          1222
5 2013     1     1     1505            1310      115      1638          1431
```

```
6 2013 1 1 1525 1340 105 1831 1626
7 2013 1 1 1549 1445 64 1912 1656
8 2013 1 1 1558 1359 119 1718 1515
9 2013 1 1 1732 1630 62 2028 1825
10 2013 1 1 1803 1620 103 2008 1750
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>
```

```
count(delay_2)
```

```
# A tibble: 1 x 1
  n
  <int>
1 10200
```

```
flights %>%
  filter( carrier == "DL" | carrier == "AA" | carrier == "UA") -> carriers
head(carriers, 10)
```

```
# A tibble: 10 x 19
  year month day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int> <int> <int> <dbl> <int> <int>
1 2013 1 1 517 515 2 830 819
2 2013 1 1 533 529 4 850 830
3 2013 1 1 542 540 2 923 850
4 2013 1 1 554 600 -6 812 837
5 2013 1 1 554 558 -4 740 728
6 2013 1 1 558 600 -2 753 745
7 2013 1 1 558 600 -2 924 917
8 2013 1 1 558 600 -2 923 937
9 2013 1 1 559 600 -1 941 910
10 2013 1 1 559 600 -1 854 902
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>
```

```
count(carriers)
```

```
# A tibble: 1 x 1
```

```

n
<int>
1 139504

```

2. I could not get the ‘and’ syntax to work so I was not able to complete this in one pipe.

```

flights %>%
  filter(arr_delay >= 120, dep_delay <= 0, between(month, 7, 8)) -> temp_flights
flights %>%
  filter(arr_delay >= 120, dep_delay <= 0, month == 6, between(day, 21, 30)) -> temp_flights2
flights %>%
  filter(arr_delay >= 120, dep_delay <= 0, month == 9, between(day, 1, 22)) -> temp_flights3

aflights <- rbind(temp_flights, temp_flights2)
flights2 <- rbind(aflights, temp_flights3)
flights2 %>%
  arrange(month)

```

```

# A tibble: 16 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>    <int>        <int>     <dbl>    <int>        <int>
1 2013     6    24    1602        1605      -3    2134        1916
2 2013     6    27    2052        2100      -8     13        2210
3 2013     6    30    1423        1425      -2    1816        1554
4 2013     7     1     905        905       0    1443        1223
5 2013     7     7    1659        1700      -1    2050        1823
6 2013     7     7    1727        1730      -3    2203        1951
7 2013     7     7    1746        1755      -9    2133        1921
8 2013     7     7    1823        1830      -7    2201        1955
9 2013     7    22    1555        1600      -5    2139        1938
10 2013    7    22    1606        1615      -9    2056        1831
11 2013    7    22    1628        1630      -2    2151        1939
12 2013    7    28    1710        1711      -1    2248        2039
13 2013    8     8    1457        1500      -3    1828        1624
14 2013    8    13     657        659      -2    1015        814
15 2013    8    28    1157        1200      -3    1520        1316
16 2013    9    19     656        700      -4    1037        833
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>

```

```
head(flights2, 10)

# A tibble: 10 x 19
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>          <int>     <dbl>     <int>          <int>
1 2013     7     1      905            905       0    1443        1223
2 2013     7     7     1659           1700      -1    2050        1823
3 2013     7     7     1727           1730      -3    2203        1951
4 2013     7     7     1746           1755      -9    2133        1921
5 2013     7     7     1823           1830      -7    2201        1955
6 2013     7    22     1555           1600      -5    2139        1938
7 2013     7    22     1606           1615      -9    2056        1831
8 2013     7    22     1628           1630      -2    2151        1939
9 2013     7    28     1710           1711      -1    2248        2039
10 2013    8     8     1457           1500      -3    1828        1624
# i 11 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
# tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
# hour <dbl>, minute <dbl>, time_hour <dttm>
```

3. The destinations that are the shortest distance are La Guardia (LGA) and Philadelphia International (PHL).

```
flights %>%
  select(distance, dest) %>%
  arrange(distance) %>%
  head(2)
```

```
# A tibble: 2 x 2
  distance dest
  <dbl> <chr>
1      17 LGA
2      80 PHL
```

```
data(airports)
```

4. There are some flights where their scheduled time and delay times are not equal to their departure times. These could be due to inaccuracies when inputting data or a complication with combining a negative delay time with the scheduled departure time.

```

flights %>%
  mutate(dep_time_true = dep_time == dep_delay + sched_dep_time) %>%
  filter(dep_time_true == FALSE)

# A tibble: 99,777 x 20
  year month   day dep_time sched_dep_time dep_delay arr_time sched_arr_time
  <int> <int> <int>     <int>        <int>      <dbl>    <int>        <int>
1 2013     1     1      554          600       -6     812        837
2 2013     1     1      555          600       -5     913        854
3 2013     1     1      557          600       -3     709        723
4 2013     1     1      557          600       -3     838        846
5 2013     1     1      558          600       -2     753        745
6 2013     1     1      558          600       -2     849        851
7 2013     1     1      558          600       -2     853        856
8 2013     1     1      558          600       -2     924        917
9 2013     1     1      558          600       -2     923        937
10 2013    1     1      559          600       -1     941        910
# i 99,767 more rows
# i 12 more variables: arr_delay <dbl>, carrier <chr>, flight <int>,
#   tailnum <chr>, origin <chr>, dest <chr>, air_time <dbl>, distance <dbl>,
#   hour <dbl>, minute <dbl>, time_hour <dttm>, dep_time_true <lgl>

```

- Frontier Airlines has the worst average departure time with a mean of over 20 minutes.

```

flights[!is.na(flights$dep_delay),] %>%
  summarize(.by = carrier, avg_delay = mean(dep_delay)) %>%
  arrange(desc(avg_delay))

```

```

# A tibble: 16 x 2
  carrier avg_delay
  <chr>      <dbl>
1 F9          20.2
2 EV          20.0
3 YV          19.0
4 FL          18.7
5 WN          17.7
6 9E          16.7
7 B6          13.0
8 VX          12.9
9 OO          12.6
10 UA         12.1

```

11	MQ	10.6
12	DL	9.26
13	AA	8.59
14	AS	5.80
15	HA	4.90
16	US	3.78

```
data(airlines)
```

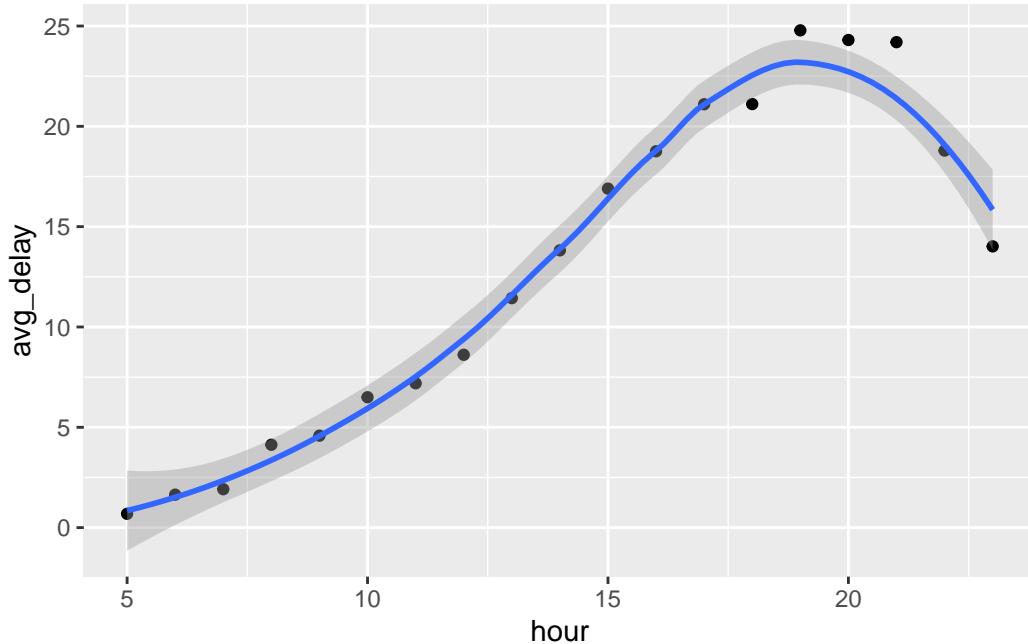
6. c. You should schedule your flights for 5 am to minimize the expected delay time.

```
library(ggplot2)

flights[!is.na(flights$dep_delay),] %>%
  summarize(.by = hour, avg_delay = mean(dep_delay)) %>%
  arrange(desc(avg_delay)) -> flights_delay

ggplot(flights_delay, aes(hour, avg_delay))+
  geom_point()+
  geom_smooth()

`geom_smooth()` using method = 'loess' and formula = 'y ~ x'
```



7. The result surprises me because I would have thought that the select function, like the rest of R, was case sensitive, but it is not. The default setting is for the helper functions to be case sensitive, you can override this by using ignore.case.

```
flights |> select(contains("TIME"))
```

```
# A tibble: 336,776 x 6
  dep_time sched_dep_time arr_time sched_arr_time air_time time_hour
  <int>        <int>    <int>        <int>     <dbl> <dttm>
1      517          515     830          819      227 2013-01-01 05:00:00
2      533          529     850          830      227 2013-01-01 05:00:00
3      542          540     923          850      160 2013-01-01 05:00:00
4      544          545    1004          1022     183 2013-01-01 05:00:00
5      554          600     812          837      116 2013-01-01 06:00:00
6      554          558     740          728      150 2013-01-01 05:00:00
7      555          600     913          854      158 2013-01-01 06:00:00
8      557          600     709          723       53 2013-01-01 06:00:00
9      557          600     838          846      140 2013-01-01 06:00:00
10     558          600     753          745      138 2013-01-01 06:00:00
# i 336,766 more rows
```